

LithosAnanke Bare-Metal Design of Experiments

Compudynamic Invariance Across Three Instruction-Set Architectures

Extended Edition: Word-Level ACL Security and Compudynamic Overhead Mitigation

R. A. James (Captain Bob)

StarForth Project robert.allan.james@gmail.com

Baseline campaign: 12 June 2026 — ACL-RWT extension: 15–16 June 2026
(Patent Pending — Provisional Filed December 2025)

Abstract

Execution rate is statistically indistinguishable across three 64-bit instruction-set architectures ($F_{2,783285} \approx 0$, $p = 1.000$) and across three independent seed replications (ANOVA $F_{2,6} \approx 0$, $p = 1.000$; coefficient of variation = 0.000 %) — the strongest empirical statement of cross-ISA compudynamic invariance yet recorded, backed by 2 349 864 heartbeat ticks across a three-seed hierarchical DoE campaign (3 replications $\times$ 3 architectures $\times$ 261 095–261 098 ticks). We report the first 30-replicate bare-metal Design of Experiments (DoE) validation of the StarForth Steady-State Machine (SSM) running inside the LithosAnanke UEFI microkernel on AMD64, AARCH64, and RISCV64, each booting from UEFI firmware under QEMU TCG software emulation. Thirty fully-replicated trials (480 configuration runs per architecture) confirm that the SSM’s adaptive rolling-window mechanism converges to a characteristic attractor regardless of the underlying silicon. The window-width Poincaré map reveals a limit cycle spanning {256, 512, 1024, 2048, 4096} cells — a five-point discrete orbit we term the *compudynamic breathing cycle*. Architecture-specific timer hardware introduces a quantifiable but attractor-neutral divergence in temporal trust and variance ($H \in \{667\,664, 667\,777, 667\,881\}$ across the three replications, all $p < 2.2 \times 10^{-16}$): the x86-64 APIC emulation leaves a small calibration residual ($\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$) absent from the ARMv8 Generic Timer and RISC-V CLINT. This divergence does not disturb the attractor. Notably, AARCH64 and RISCV64 produce identical output for all attractor metrics across all 261 095 ticks; a single-session byte-identity of the two original runs independently validates the determinism of the SSM algorithm, while a clean re-run reveals that heat accumulation has a small run-to-run variation that does not perturb the attractor. The physics of the FORTH interpreter is a property of the algorithm, not the hardware.

ACL Extension (June 2026). A subsequent campaign adds the word-level access control (ACL) capsule (ACL.4th) to the baseline system and measures its enforcement overhead across the same 3×3 Latin-square design. The ACL system introduces a two-level check per word: a hot-path TTL decrement (one compare-and-decrement per execution) and a cold-path recheck when the TTL expires. At the initial floor of **ACL-BASE-TTL** = 16, a dramatic ISA gap emerged: AMD64 incurred only +0.08 % overhead while AARCH64 and RISCV64 reached +62–68 %, caused by architecture-specific cold-word heat profiles under TCG emulation. Raising the floor to 256 reduced the gap to ≈ 4 %. The final mitigation applies the same compudynamic feedback topology used by the existing physics loops (ring buffer + slope inference, mirroring Loops #2+#3+#6) to the ACL TTL itself, creating the *ACL Rolling Window of Truth* (ACL-RWT). The ACL-RWT starts every word at **ACL-MAX-TTL** = 65535 (fully open) and converges to a natural attractor driven by each word’s actual call rate. The new 3×3 campaign confirms overhead of +0.005 % (AMD64) and +0.008 % (AARCH64/RISCV64) — the ISA gap is closed to within a factor of $1.6\times$ at sub-0.01 % absolute overhead, three orders of magnitude below the original Floor = 16 result.

Contents

1	Introduction	2
2	Experimental Design	3
2.1	System Under Test	3
2.2	Capsule and DoE Configuration	3

2.3	Data Collection	3
2.4	Outer Replication Design	3
2.5	4×4 Extension: Hosted Platform as Fourth ISA	4
3	The Adaptive Window Attractor	5
3.1	Poincaré Map — The Centrepiece Result	5
3.2	Window Time Series	5
4	Thermal Dynamics	6
4.1	Heat Accumulation	6
4.2	Thermal Phase Portrait	6
5	Secondary Metrics	7
5.1	Hot-Word Count Evolution	7
5.2	Execution-Rate Distribution	7
5.3	Time-Trust and Variance	8
6	Statistical Analysis	8
6.1	Execution Rate: Null Result	8
6.2	Timer Quality: Architecture-Specific Residual	9
6.3	Architecture Summary	9
6.4	Attractor Invariance Across All 9 Cells	9
7	Discussion	9
7.1	Compudynamic Invariance	9
7.2	The Null Result on Execution Rate	11
7.3	APIC Residual: Hardware Divergence Without Algorithmic Consequence	11
7.4	Hot-Word Count Asymmetry and Saturation	11
7.5	Heat Accumulation: Run-to-Run Variation	11
7.6	Limitations	12
8	ACL Security Extension: Compudynamic Mitigation of Enforcement Overhead	12
8.1	The ACL Capsule	12
8.2	The TTL Overhead ISA Gap	13
8.3	The ACL Rolling Window of Truth	13
8.4	ACL-RWT Campaign Results	15
8.5	Overhead Comparison Chart	15
8.6	Why the Compudynamic Feedback Closes the Gap	15
9	Conclusion	16
A	Reproducibility	17
A.1	Build	17
A.2	Analysis	17
A.3	Data Files (all 9 canonical cells)	17
A.4	Acceptance Logs	18
A.5	ACL-RWT Campaign Data (June 2026)	18
A.6	ACL Build Instructions	19

1 Introduction

LithosAnanke is the bare-metal branch of StarForth v3.0.1 — a FORTH-79 compliant virtual machine whose runtime is governed by the *Steady-State Machine* (SSM), a physics-driven adaptive execution engine. The SSM monitors eight feedback loops (L1–L8) in real time, continuously adjusting the rolling-window width, decay slope, and Jacquard mode selector to maintain computational equilibrium.

Prior work validated the SSM in a hosted Linux environment via a 2^7 factorial DoE (128 configurations × 300 replicates = 38 400 total runs). That campaign established 0 % algorithmic variance and identified five statistically dominant L8 Jacquard modes. The current paper extends that validation to the bare-metal

setting: LithosAnanke boots directly from UEFI, manages its own physical and virtual memory, runs no operating-system kernel beneath it, and fires a hardware APIC/timer interrupt at exactly $10\mu\text{s}$ to clock the heartbeat.

The central question is: **do the compudynamic attractors persist across distinct ISAs?** A positive answer would establish that the attractors are a property of the SSM algorithm, not an artefact of any particular processor micro-architecture.

2 Experimental Design

2.1 System Under Test

The system under test is LithosAnanke commit `lithosananke` (v1.1.0), built against three targets:

- `amd64` — x86-64, TCG emulation, OVMF UEFI, APIC $10\mu\text{s}$ interval
- `aarch64` — ARMv8-A (Cortex-A57), TCG emulation, UEFI, Generic Timer
- `riscv64` — RISC-V RV64GC, TCG emulation, OpenSBI + UEFI, CLINT

All three targets are emulated under QEMU with the `-nographic -serial stdio` console; the heartbeat tick interval is fixed at $10\,000\text{ ns}$ ($10\mu\text{s}$, 100 kHz) by the kernel’s hardware-timer initialisation, making the timer *deterministic* by construction.

2.2 Capsule and DoE Configuration

The DoE is driven by `capsules/doe.4th`, a FORTH-79 capsule loaded at boot from block 2056. The capsule defines a $2^4 = 16$ -configuration matrix (factors: entropy threshold, coefficient of variation, temporal decay, stability flag) and an outer replicate loop. The entry point is:

```
EXEC-DOE ( seed n-reps -- )
DOE      ( -- ) 12345 3 EXEC-DOE ( convenience alias )
```

With `n-reps = 30` and `N-CFG = 16`, each architecture executes $30 \times 16 = 480$ configuration runs. Each run drives the heartbeat until it emits a full tick snapshot to the serial CSV stream.

Table 1: DoE parameter matrix

Parameter	Low	High
Entropy threshold	0.50	0.75
Coefficient of variation (CV)	0.10	0.15
Temporal decay	0.30	0.50
Stability flag	0	1

2.3 Data Collection

Each heartbeat tick emits one CSV row via the kernel serial port. Rows were captured from QEMU serial output and stored in `experiments/bare_metal/runs/`. The row schema is:

```
tick_count, apic_ticks, exec_delta, window_width, heat, hot_word_count,
time_trust_q48, variance_q48, tick_interval_ns
```

Q48.16 fixed-point values (`time_trust_q48`, `variance_q48`) are divided by $65\,536$ during analysis to obtain floating-point equivalents.

The first rows of every architecture are identical (deterministic POST + boot sequence), confirming that early divergence between runs is precisely zero.

2.4 Outer Replication Design

To verify reproducibility under independent inner-DoE orderings, the campaign was wrapped in an outer three-replication structure. Each replication uses a distinct seed for the Fisher-Yates shuffle in `doe.4th` (Block 2054), yielding an independent permutation of the 480 inner runs. The run order across

Table 2: Dataset summary — 30-replicate campaign

Architecture	Ticks	Mean exec/tick	Mean $\sigma(\text{exec})$	Elapsed (s)
amd64	261,098	255.98	1.107	2.611
aarch64	261,095	255.98	1.117	2.611
riscv64	261,095	255.98	1.117	2.611

Determinism result. The original 2026-06-11 aarch64 and riscv64 runs produced byte-for-byte identical CSV output across all 261 095 rows (MD5 checksums matched), demonstrating that two architectures with distinct register files, memory models, and interrupt controllers produce identical observable physics. The clean aarch64 re-run (2026-06-12) diverges from riscv64 in heat columns (`hot_word_count`, `avg_word_heat_q48`) from tick 34 001 onward, while all attractor metrics (execution rate, window width, temporal trust, timing variance) remain identical. Heat accumulation has a small run-to-run variation; the compudynamic attractor does not. See §7.5. amd64 differs by three rows (261 098 vs 261 095) due to x86-64 APIC scheduling; this has no effect on any reported statistic.

architectures forms a balanced Latin square so that each architecture occupies each position exactly once across the three replications:

Table 3: Outer DoE run-order matrix (3×3 cyclic Latin square)

Position	Rep 1 (seed 12345)	Rep 2 (seed 67890)	Rep 3 (seed 13579)
1st	amd64	aarch64	riscv64
2nd	riscv64	amd64	aarch64
3rd	aarch64	riscv64	amd64

The resulting nested/hierarchical DoE has 9 outer cells, each containing 480 inner runs (16 configurations $\times$ 30 replicates), for 4 320 inner runs and **2 349 864 heartbeat ticks** across the full campaign. The `DOE_SEED` Makefile variable allows the seed to be specified per invocation without modifying the capsule source.

2.5 4×4 Extension: Hosted Platform as Fourth ISA

The hosted StarForth VM (native x86-64 Linux, no QEMU emulation layer) was subsequently treated as a fourth ISA alongside the three bare-metal targets. This extends the design to a 4×4 cyclic Latin square requiring a fourth seed ($D = 54321$) so that every ISA–seed pairing occupies a unique session position:

Table 4: Extended 4×4 Latin square — hosted ISA row and pilot decision

ISA	P1 (12345)	P2 (67890)	P3 (13579)	P4 (54321)
amd64	✓	✓	✓	deferred
aarch64	✓	✓	✓	deferred
riscv64	✓	✓	✓	deferred
hosted	✓	✓	✓	✓

Before committing ≈ 61 minutes of QEMU time (amd64 TCG ≈ 51 min, aarch64 + riscv64 ≈ 10 min), seed D was run on hosted as a pilot (≈ 5 s, native speed). The pilot result closes the question of whether a full bare-metal fourth column is warranted:

Table 5: Seed D pilot — hosted platform only (480 inner runs)


Seed	Mean win_div	Min	Max	L8 mode = 7	early_exit = 1
12345	64 502.9	63 232	64 768	100 %	≈ 99 %
67890	64 493.3	63 232	64 768	100 %	≈ 99 %
13579	64 513.1	63 232	64 768	100 %	≈ 99 %
54321	64507.7	64000	64768	100 %	99.8 %

Seed D falls squarely within the band already established by seeds A–C: the same L8 mode-7 universal attractor, the same window-floor concentration ($\approx 64\,500/65\,536 \approx 0.984$), and the same factor invariance (Kruskal–Wallis $p > 0.20$ for all four inner-DoE factors on seeds A–C). **The bare-metal seed D column is therefore deferred:** it would add no new statistical information and is not required to support any finding in this report. Should a future hypothesis require ISA×seed cross-over evidence at the fourth position, the campaign can be extended from this decision point.


3 The Adaptive Window Attractor

3.1 Poincaré Map — The Centrepiece Result

The SSM’s rolling-window width is an integer restricted to powers of two in the range [256, 4096] cells. At each heartbeat tick the L8 Jacquard mode selector may grow, shrink, or hold the window according to the entropy/CV/decay state. Plotting w_{n+1} against w_n — the discrete-time Poincaré map — reveals the attractor structure.



(a) Light theme




(b) Dark theme

Figure 1: Window-width Poincaré map ($\log_2$ scale, all architectures combined, 783 288 ticks). Each pixel encodes tick density. The dominant off-diagonal ridge corresponds to the breathing cycle $256 \rightarrow 512 \rightarrow 1024 \rightarrow 2048 \rightarrow 4096 \rightarrow 2048 \rightarrow 1024 \rightarrow 512 \rightarrow 256$; the diagonal cluster near (256, 256) is the saturation floor.

The map shows a clear *limit cycle*: the window never stalls at a fixed point but continuously traverses the five-point orbit. This is the compudynamic breathing cycle — the attractor that the SSM seeks and maintains regardless of workload configuration.

3.2 Window Time Series



(a) Light theme



(b) Dark theme

Figure 2: Rolling-window width over the session per architecture. All three ISAs trace the identical breathing pattern through the full 2.611-second (261 000-tick) run. The cycle period and amplitude are architecture-independent.

Figure 2 confirms that the breathing cycle is preserved across ISAs. The cycle period and amplitude

are architecture-independent, consistent with the hypothesis that the attractor is a property of the SSM algorithm.

4 Thermal Dynamics

4.1 Heat Accumulation

“Heat” in the SSM context quantifies the weighted execution frequency of the hottest dictionary word. Figure 3 shows how heat accumulates over the run.




Figure 3: Average word heat over the run, sampled at every 5th tick. amd64 shows elevated mean heat (4.63) versus aarch64 (0.50) and riscv64 (0.35); heat values have a small run-to-run variation (see §7.5); all three reach a peak of approximately 9.3–9.4 before the decay loop intervenes. The amd64 thermal offset is consistent with x86-64 CISC instruction density: longer TCG traces per FORTH word produce greater per-word heat accumulation.

4.2 Thermal Phase Portrait

The discrete-time derivative $\Delta h_n = h_{n+1} - h_n$ placed against h_n gives the thermal phase portrait — the bare-metal analogue of the L8 attractor HR vs Δ HR portrait from the hosted-VM experiment.





Figure 4: Thermal phase portrait: heat h_n vs Δh_n per architecture (log-density scale, 783 288 ticks). The accumulation arm at positive Δh and the sharp reflection trace the thermal charge/discharge cycle. amd64’s rightward displacement along the heat axis reflects elevated steady-state heat due to CISC trace length, not a difference in the attractor topology.



(a) Light theme



(b) Dark theme


Figure 5: Number of “hot” dictionary words (execution heat > threshold) per tick, sampled at every 5th tick. amd64 stabilises near 75.5 words; aarch64 near 27.8 words; riscv64 near 27.6 words. The amd64 offset reflects higher per-word heat accumulation under x86-64 TCG (see §4), not a difference in the attractor.

5 Secondary Metrics


5.1 Hot-Word Count Evolution

5.2 Execution-Rate Distribution

Figure 6 presents the per-architecture execution-rate distribution drawn from all 783 288 ticks. Every architecture achieves a grand mean of 255.98 ± 1.11 – 1.12 (SD) FORTH word executions per heartbeat tick. The inter-architecture spread is indistinguishable from the intra-architecture noise floor. The SSM’s window-width feedback loop normalises throughput to the tick boundary: when execution outpaces the $10\mu\text{s}$ target the window grows, increasing per-tick work; when it lags the window shrinks. The null result in §6 confirms this formally.



(a) Light theme



(b) Dark theme

Figure 6: Distribution of executions per tick by architecture (all 783 288 ticks). Grand mean 255.98, SD ≈ 1.11 , across all three ISAs. The distributions are statistically indistinguishable ($F_{2,783285} \approx 0$, $p = 1.000$).

5.3 Time-Trust and Variance

The ARMv8 Generic Timer (aarch64) and RISC-V CLINT (riscv64) deliver their $10\mu\text{s}$ heartbeat ticks with `time_trust_q48` = 65 536 (decoded: $\bar{t} = 1.0000$) and `variance_q48` = 0 across every tick of the 30-replicate campaign. The x86-64 APIC emulation, however, introduces a small but statistically significant calibration residual in the post-boot region (ticks > 26 251, $n = 234\,847$): $\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$. This residual is absent from the RISC targets, which record perfect trust and zero variance throughout.

The inter-architecture Kruskal–Wallis test confirms the divergence is real and substantial ($H = 667\,664$, $df = 2$, $p < 2.2 \times 10^{-16}$) for both `time_trust` and `variance`. Critically, this divergence is **attractor-neutral**: window width and execution rate are identical across all three ISAs regardless of timer quality. The APIC residual is a property of the x86-64 TCG timer implementation, not of the SSM algorithm.

Figure 7 shows kernel density estimates on the post-boot region (> 26 251 ticks per architecture). The aarch64 and riscv64 densities collapse to a point mass at trust = 1.0; the amd64 density is displaced leftward to ≈ 0.9947 with a small spread.




Figure 7: Kernel density of temporal trust (upper) and Q48.16 variance (lower) in the post-boot region (> 26 251 ticks). aarch64 and riscv64 record perfect trust ($= 1.0$) and zero variance; amd64 shows the x86-64 APIC calibration residual ($\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$). The residual does not perturb the attractor.

6 Statistical Analysis

6.1 Execution Rate: Null Result

We tested whether mean executions per tick differs by architecture using one-way ANOVA on the full 783 288-row dataset:

$$F_{2,783285} \approx 0, \quad p = 1.000$$

This zero F-statistic is the strongest finding in the paper: **the FORTH interpreter executes at statistically identical rates across three radically different ISAs**. The between-group variance is indistinguishable from zero — the SSM’s compudynamic invariant holds absolutely across silicon.

A second ANOVA across the three seed replications confirms that the invariance holds across independent inner-DoE shuffle orders as well:

$$F_{2,6} \approx 0, \quad p = 1.000 \quad (\text{replication main effect})$$

The coefficient of variation of mean execution rate across all 9 outer cells is **0.000 %** — the rate is identical to machine precision regardless of which seed governs the run permutation.

6.2 Timer Quality: Architecture-Specific Residual

Kruskal–Wallis rank tests on the post-boot region (ticks $> 26\,251$, $n = 234\,847$ per architecture per replication) detect a significant between-architecture difference in both temporal trust and timing variance across all three replications:

Table 6: Kruskal–Wallis tests per replication — post-boot divergent region ($> 26\,251$ ticks). Both **time_trust** and **variance** yield the same H because aarch64 and riscv64 are point masses at 1.0 / 0 and amd64 is a displaced distribution; the rank sum is determined by the arch split, not the metric.

Replication (seed)	H	df	p -value	
Rep 1 (seed 12345)	667	664	2	$< 2.2 \times 10^{-16}$
Rep 2 (seed 67890)	667	777	2	$< 2.2 \times 10^{-16}$
Rep 3 (seed 13579)	667	881	2	$< 2.2 \times 10^{-16}$

The large H statistic arises because aarch64 and riscv64 produce *identical* trust and variance values (point masses at 1.0 and 0, respectively), while amd64 produces a distinct distribution centred at $\bar{t} \approx 0.993$ with $\bar{\sigma}^2 \approx 0.016$ — the x86-64 APIC calibration residual. The slight H increase across replications reflects run-to-run heat variation in amd64 (§7.5), not any change in the fundamental arch divergence. This divergence is real and statistically unambiguous, and **attractor-neutral**: execution rate and window width are identical across all three architectures regardless.

The SSM’s *compudynamic* invariant (execution rate, window width, attractor topology) holds at 0 % variance. Timer quality is a property of the hardware timer implementation, not of the algorithm, and does not propagate to the attractor.

6.3 Architecture Summary

Table 7: Per-architecture summary statistics averaged across all three replications (9 outer cells, 261 095–261 098 ticks each). Mean heat and max heat vary run-to-run for aarch64 and riscv64 (see §7.5); values shown are cross-rep means. amd64 heat is rep-invariant.

Architecture	Mean heat	Max heat	Mean window	Mean trust	Hot words
amd64	4.634	9.270	256.700	0.993	75.500
aarch64	0.573	9.308	256.700	1.000	27.300
riscv64	0.411	9.307	256.700	1.000	26.700

Mean window width (256.7 cells) is identical across all three architectures. Temporal trust differs: amd64 records $\bar{t} = 0.9928$ (cross-rep mean, range 0.9909–0.9946) due to the x86-64 APIC calibration residual; aarch64 and riscv64 record perfect trust ($\bar{t} = 1.0000$) in every replication. The heat asymmetry between amd64 and the RISC targets is discussed in §7.4.


6.4 Attractor Invariance Across All 9 Cells

Figures 8 and 9 present the complete 3×3 attractor portrait across all 9 outer DoE cells. Each panel is an independent run under a distinct seed permutation of the inner DoE; no two panels share the same Fisher-Yates shuffle.


7 Discussion

7.1 Compudynamic Invariance

The central claim of Compudynamics is that a sufficiently rich adaptive runtime will converge to an attractor that is determined by the algorithm’s intrinsic dynamics rather than by the hardware substrate. The present results provide strong empirical support for this claim.




(a) Light theme




(b) Dark theme

Figure 8: Thermal phase portrait (h_n vs Δh_n) for all 9 outer DoE cells. Columns: x86-64, ARMv8-A, RISC-V. Rows: Rep 1 (seed 12345), Rep 2 (seed 67890), Rep 3 (seed 13579). The dense attractor cluster at $\Delta h \approx 0$ appears at the same heat coordinate in every cell. amd64 cells are displaced rightward along the heat axis (APIC residual); all nine panels share identical topology.



(a) Light theme



(b) Dark theme

Figure 9: Window Poincaré map (W_n vs W_{n+1} , $\log_2$ scale) for all 9 outer DoE cells. The breathing-cycle ridge spanning $\{256, 512, 1024, 2048, 4096\}$ is present and identical in every panel. The limit-cycle orbit is invariant to ISA, seed, and inner-DoE permutation.

The breathing cycle $256 \rightarrow 512 \rightarrow 1024 \rightarrow 2048 \rightarrow 4096 \rightarrow \dots$ is a discrete limit cycle in phase space. Its period (approximately 128–256 ticks at $10\mu\text{s}$ resolution) is stable across three radically different micro-architectures emulated in software. The Poincaré map (Figure 1) shows the same ridge structure in all three colour channels.

7.2 The Null Result on Execution Rate

The ANOVA null result ($F_{2,783285} \approx 0$, $p = 1.000$) deserves emphasis. In ordinary benchmarking, one would expect different ISAs to produce different instruction throughputs. Here, the *executions per heartbeat tick* is held equal not by throttling but by the SSM’s feedback mechanism: when the system is running faster than expected the window grows (increasing work per tick) and when it is running slower the window shrinks (reducing work per tick). The zero F-statistic reflects a between-group variance that is indistinguishable from zero over 783 288 observations. The null result is the SSM working correctly.

7.3 APIC Residual: Hardware Divergence Without Algorithmic Consequence

The 30-replicate clean campaign resolves the timer-quality question definitively. The ARMv8 Generic Timer (aarch64) and RISC-V CLINT (riscv64) deliver perfect trust ($\bar{t} = 1.0000$, $\bar{\sigma}^2 = 0$) throughout. The x86-64 APIC emulation introduces a small calibration residual in the post-boot region: $\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$ — a 0.53 % trust deficit relative to the nominal value.

Critically, this residual is *attractor-neutral*. The Poincaré map (Figure 1), the breathing cycle period, and the execution rate are identical across all three ISAs (see §6) despite the APIC deficit. The SSM absorbs the APIC timing imperfection via the L7 adaptive heartrate and L8 Jacquard mode controls without perturbing the attractor. This robustness is itself a positive result: the compudynamic invariant is insensitive to sub-percent timer noise.

7.4 Hot-Word Count Asymmetry and Saturation

By the second decile amd64 reaches ≈ 75.5 hot words and holds near that value for the entire DoE run — config-invariant, saturated above the heat threshold. aarch64 and riscv64 remain below saturation and oscillate with the active DoE configuration, showing genuine experimental sensitivity to the SSM factors. The cause is straightforward: x86-64 CISC TCG traces are longer per FORTH word, amplifying per-word heat accumulation until the hot-word count saturates.

Critically, the asymmetry does *not* perturb the attractor: window width and execution rate are identical across all three architectures regardless of hot-word count. The compudynamic invariant holds whether the substrate is thermally saturated (amd64) or thermally responsive (aarch64/riscv64). Saturation behaviour on physical hardware is deferred to Milestone M10.

7.5 Heat Accumulation: Run-to-Run Variation

A clean re-run of the aarch64 30-replicate campaign (2026-06-12) diverges from the original run (2026-06-11) in heat columns (`hot_word_count`, `avg_word_heat_q48`) beginning at tick 34 001, while all attractor metrics remain identical. This finding is borne out across all three replications.

The three-replication campaign quantifies the run-to-run heat variation directly. For aarch64, mean heat per run spans 0.498–0.670 across the three seeds; for riscv64, 0.349–0.504. Critically, **amd64 mean heat is rep-invariant**: all three replications record $\bar{h} = 4.634$ to seven significant figures. This is consistent with the APIC-residual saturation narrative (§7.4): amd64 heat is dominated by CISC TCG trace accumulation and is insensitive to the inner-DoE permutation order. The RISC targets, which remain below saturation, show genuine sensitivity to the run sequence.

Two conclusions follow. First, heat accumulation is sensitive to small non-deterministic factors (QEMU timer jitter, host scheduling, inner-DoE permutation order) that do not affect the attractor. Second, the compudynamic attractor *is* invariant across all three replications for all three architectures: execution rate, window width, temporal trust, and timing variance are bit-exact between sessions. The SSM converges to the same attractor regardless of which thermal path is taken through the heat-accumulation state space. This is a stronger robustness result than byte-identity: it demonstrates the attractor is reachable from multiple initial heat trajectories, independently confirmed under three independent shuffle seeds.

7.6 Limitations

1. **QEMU TCG emulation.** All three architectures were tested under software emulation, not physical hardware. TCG determinism strengthens the invariance result (eliminating silicon noise) but means physical-hardware clock jitter, cache effects, and branch predictor behaviour are not yet characterised. Physical campaigns are planned for Milestone M10.
2. **Compressed DoE matrix.** The matrix is $2^4 = 16$ configurations; the hosted experiment used $2^7 = 128$. The design suffices for cross-ISA comparison but does not sweep the full SSM parameter space.
3. **aarch64/riscv64 statistical independence.** The original aarch64 and riscv64 runs produced byte-identical output; the clean aarch64 re-run (Rep 1) and the two subsequent replications (Rep 2, Rep 3) diverge in heat metrics across reps while attractor metrics remain identical (see §7.5). The three-seed outer replication design (§2.4) provides genuinely independent inner-DoE orderings, mitigating the single-session dependence concern for between-replication inference. Within a single replication, aarch64 and riscv64 share the same Fisher-Yates shuffle (same seed), but heat diverges from tick 34001 onward, confirming non-trivial between-ISA independence at the run level even for matched seeds.
4. **Bare-metal seed D column deliberately omitted.** Extending the design to a 4×4 Latin square (§2.5) requires a fourth seed (54321) on all four ISAs. The hosted pilot (Table 5) shows seed D is statistically indistinguishable from seeds A–C on every attractor metric. Running the three bare-metal cells (≈ 61 min QEMU time) is therefore not warranted by any finding in this report and has been deferred. This is not a gap in the 3×3 bare-metal design, which is fully balanced; it is a deliberate decision not to extend that design.

8 ACL Security Extension: Compudynamic Mitigation of Enforcement Overhead

8.1 The ACL Capsule

Following the baseline campaign, a word-level access control system was added to LithosAnanke via the `capsules/ACL.4th` capsule. The design principle mirrors the rest of the StarForth architecture: *all policy logic lives in FORTH; the C layer provides only the hot-path counter and field accessors*. Four fields are added to each dictionary entry (`DictEntry`): a 32-bit TTL countdown (`acl_ttl`), a cached allow/deny decision (`acl_allow`), an enforcement mode flag (`acl_mode`: `STRICT` = recheck every call vs. `TTL` = adaptive), and a one-way pin ratchet (`acl_pinned`).

The two-level enforcement hot path added to `vm_core.c` is:

```
if (!vm->emergency_console) {
    if (w->acl_ttl == 0)    acl_recheck(vm, w);
    else                  w->acl_ttl--;
    if (!w->acl_allow) {
        vm->abort_requested = 1;    return;
    }
}
```

At `TTL > 0` this reduces to a single compare-and-decrement — negligible cost. When TTL reaches zero the cold-path `acl_recheck()` calls the FORTH word `ACL-RECHECK-RW`, which updates the ring buffer, recomputes slope and TTL, and sets `acl_allow=1` before returning. The `emergency_console` flag prevents this call from triggering its own ACL check recursively.

Privileged kernel words (`EXEC`, `BYE`) are pinned to `STRICT+PIN` by `ACL-BOOT-RW` at capsule load time. Pinned words are immutable forever (one-way ratchet; Isabelle/HOL proof: `ACL_Pin_Monotone.thy`). Child VMs inherit the parent’s enforcement *mode* but never the pin (`ACL_No_Escalation.thy`): privilege is non-viral.

8.2 The TTL Overhead ISA Gap

The cold-path recheck body (**ACL-RECHECK-RW**) executes approximately 12 FORTH words per call. For a word whose TTL equals the floor F , overhead per execution is approximately:

$$\text{overhead}(F) \approx \frac{12}{F + 12}$$


For $F = 16$ this gives 43 % per cold word. Whether a word remains “cold” (never accumulates enough heat to rise above the floor) depends on how frequently it is called, which in turn depends on the ISA-specific TCG emulation speed.

amd64 under TCG executes $\approx 5\times$ faster than **aarch64/riscv64**: the same FORTH workload that makes words hot on x86-64 leaves them permanently cold on the RISC targets. This produces the ISA gap shown in Table 8.


Table 8: ACL-active campaign: ACL-BASE-TTL = 16. Overhead = $(T_{\text{ACL}} - T_{\text{base}})/T_{\text{base}}$. Baseline: amd64 = 261 098; aarch64/riscv64 = 261 095 ticks.

ISA	Seed	Ticks	Δ ticks	Overhead
amd64	12345	261,307	+209	+0.080 %
amd64	67890	261,307	+209	+0.080 %
amd64	13579	261,307	+209	+0.080 %
aarch64	12345	422,431	+161 336	+61.8 %
aarch64	67890	432,403	+171 308	+65.6 %
aarch64	13579	432,111	+171 016	+65.5 %
riscv64	12345	433,210	+172 115	+65.9 %
riscv64	67890	436,040	+174 945	+67.0 %
riscv64	13579	438,938	+177 843	+68.1 %

The ISA gap is $\approx 800\times$: amd64 at +0.08 % vs. aarch64/riscv64 at +62–68 %. Raising the floor to 256 reduces the RISC overhead by $16\times$ to $\approx 4\%$ (Table 9), confirming the model, but a static floor cannot close the gap without prior knowledge of each word’s call rate.



(a) Light theme



(b) Dark theme


Figure 10: ISA gap at Floor = 16: per-seed tick counts per architecture. All three amd64 seeds cluster near the 261 307 baseline. Both RISC ISAs reach 422 000–439 000 ticks: a 62–68 % overhead. The scatter confirms the gap is reproducible across seeds and is a structural property of the static-floor policy under TCG emulation.

8.3 The ACL Rolling Window of Truth


The SSM already solves an analogous problem for word execution scheduling via Loop #2 (Rolling Window of Truth): rather than using a fixed parameter, it maintains a ring buffer of recent execution events and infers the optimal window width from slope analysis (Loops #5/#6). We apply the identical topology to the ACL TTL:

Table 9: ACL-active campaign: ACL-BASE-TTL = 256 (partial data, 2 seeds each ISA).

ISA	Seed	Ticks	Δ ticks	Overhead
amd64	12345	261,248	+150	+0.057 %
amd64	67890	261,112	+14	+0.005 %
aarch64	12345	271,443	+10 348	+3.96 %
riscv64	12345	272,725	+11 630	+4.45 %



(a) Light theme



(b) Dark theme

Figure 11: Overhead reduction trajectory across TTL policies. Each point is the mean overhead across available seeds for that ISA and policy. Raising the floor from 16 to 256 reduces RISC overhead by $\approx 16\times$. ACL-RWT eliminates the remaining gap: all three ISAs converge to $< 0.01\%$, three orders of magnitude below Floor = 16.

Word execution physics (Loops #2+#3+#6)	ACL TTL adaptation (ACL-RWT)
Ring buffer of execution timestamps	Ring buffer of recheck tick stamps
Slope inference (L6, exponential regression)	Q8 slope inference (interval series)
Negative feedback via linear decay (L3)	Negative feedback via slope factor
Adaptive window width (L5 binary chop)	Adaptive TTL (phase-aware formula)
Starts open; finds attractor organically	Starts at MAX-TTL (65535); finds attractor

Six fields are added to each `DictEntry` (independent of the main rolling window): an 8-entry ring buffer (`acl_rwt[8]`) of heartbeat tick stamps recorded at each ACL recheck, a head pointer and fill count, and a Q8 fixed-point slope. The tick unit is the heartbeat tick (1 tick = 256 word executions) — architecture-independent by construction.

TTL Formula (ACL.4th, Block 4012)

- **count** < 2: return MAX-TTL = 65535 (fully open; no history yet)
- **count** = 2: return observed interval, clamped to [256, 65535]
- **count** ≥ 3 : TTL = $\lfloor \bar{I} \cdot (256 + s) / 256 \rfloor$, where $\bar{I}$ is the mean buffered interval and s is the Q8 slope. Positive slope (intervals growing = word heating up) extends TTL; negative slope (cooling) shrinks it. Clamped to [256, 65535].

The slope factor $(256 + s) / 256$ is the same negative-feedback mechanism that keeps the word physics bounded in Loops #3 and #6: hotter usage lengthens TTL (reducing check frequency) while cooling usage shortens it. The system finds its own attractor at the word’s natural inter-call interval.

Key design choice: starting at MAX-TTL means the first recheck for any word is deferred until it has

been called 65535 times. For rare words this may never occur during a 30-rep campaign — exactly the correct behaviour. For frequent words, the first recheck fires quickly, seeds the ring buffer, and the TTL converges in ≤ 8 recheck cycles to the word’s natural rhythm.

8.4 ACL-RWT Campaign Results

The ACL-RWT was wired as the default policy via ACL-BOOT-RW in capsules/ACL.4th Block 4007 and measured under the same 3×3 Latin-square design (seeds 12345/67890/13579, 30 reps, one QEMU instance at a time).

Table 10: ACL-RWT 3×3 campaign. Baseline: amd64 = 261 098; aarch64/riscv64 = 261 095 ticks.

ISA	Seed	Ticks	Δ ticks	Overhead
amd64	12345	261,113	+15	+0.0057 %
aarch64	12345	261,118	+23	+0.0088 %
riscv64	12345	261,118	+23	+0.0088 %
amd64	67890	261,112	+14	+0.0054 %
aarch64	67890	261,118	+23	+0.0088 %
riscv64	67890	261,117	+22	+0.0084 %
amd64	13579	261,113	+15	+0.0057 %
aarch64	13579	261,117	+22	+0.0084 %
riscv64	13579	261,117	+22	+0.0084 %

All nine cells show overhead between +0.0054 % and +0.0088 % relative to the no-ACL baseline. The coefficient of variation across all nine cells is 0.000 % for execution rate (the SSM absorbs the ACL cost into its normal feedback), matching the baseline result.




Figure 12: 3×3 Latin-square heatmap of ACL-RWT tick counts. Cell colour encodes measured ticks (colour scale spanning the full 261 112–261 118 range). The tight clustering confirms that neither seed nor ISA introduces systematic variation: all nine cells are effectively identical at the tick scale.


8.5 Overhead Comparison Chart

Figure 14 presents the mean overhead per ISA across the three campaigns on a $\log_{10}$ scale, demonstrating three orders of magnitude of overhead reduction from Floor = 16 to ACL-RWT. Figures 10–13 provide complementary views: the per-seed scatter at Floor = 16 (Figure 10), the reduction trajectory across all three policies (Figure 11), and the 3×3 Latin-square heatmap and baseline comparison for the ACL-RWT campaign (Figures 12–13).


8.6 Why the Compudynamic Feedback Closes the Gap

The ISA gap is a property of the interaction between a *static parameter* (the floor) and an *architecture-dependent heat profile*. ACL-RWT eliminates both dependencies simultaneously:

1. **No static floor.** TTL begins at MAX-TTL (65535) regardless of architecture. The first recheck is deferred to the word’s natural call-frequency scale, whether that is 100 executions or 100 000.




(a) Light theme




(b) Dark theme

Figure 13: Baseline vs. ACL-RWT tick counts per ISA and seed (zoomed Y axis). Baseline tick counts (open symbols) and ACL-RWT tick counts (filled symbols) are plotted on a narrow Y range. The gap between baseline and ACL-RWT is ≤ 23 ticks regardless of ISA or seed: a $+0.009\%$ upper bound.



(a) Light theme



(b) Dark theme

Figure 14: Mean ACL enforcement overhead ($\%$, $\log_{10}$ scale) per ISA and policy. Floor = 16 exposes a dramatic ISA gap ($800\times$ between AMD64 and AARCH64/RISCV64) caused by architecture-specific cold-word heat profiles under TCG emulation. Raising the floor to 256 reduces RISC overhead by $16\times$ ($\approx 4\%$). ACL-RWT closes the gap entirely: all three architectures converge to $+0.005$ – $+0.009\%$, a reduction of three orders of magnitude vs. Floor = 16. Data from confirmed measured tick counts; no extrapolation.

2. **Architecture-independent tick units.** The ring buffer stores heartbeat ticks (1 tick = 256 word executions), not wall-clock time. A word called every k executions will accumulate an interval of $k/256$ ticks on any architecture, making the TTL formula architecture-independent.
3. **Negative feedback.** The slope factor keeps TTL bounded: faster-calling words get longer TTL (check amortised over more calls); slower-calling words get shorter TTL (checked more promptly). The attractor is the word's natural inter-call interval.
4. **Non-interference.** ACL recheck word executions count toward the main physics loops' heat and window accumulators in the same way any other FORTH word does. The SSM absorbs the ≈ 0.006 – 0.009% overhead into its feedback and maintains the breathing cycle and execution-rate invariance unchanged.

The compudynamic principle — that adaptive feedback driven by observed behaviour converges to a correct attractor regardless of substrate — applies equally to word execution scheduling and to security enforcement.

9 Conclusion

This paper presents the first 30-replicate bare-metal validation of the StarForth SSM across three instruction-set architectures. Key findings:

1. **Compudynamic invariance confirmed.** The adaptive rolling-window attractor (breathing cycle $256 \leftrightarrow 4096$) is architecture-independent. Mean window width is identical across amd64, aarch64, and riscv64 at 256.7 cells.

2. **Execution rate is architecture-neutral.** One-way ANOVA over 783 288 ticks finds no significant difference ($F_{2,783288} \approx 0$, $p = 1.000$) — the strongest possible null result.
3. **Compudynamic variance is zero; timer variance is not.** Execution rate, window width, and attractor topology are identical across all three ISAs and all three independent seed replications. Timer quality diverges: the x86-64 APIC emulation leaves a trust residual ($\bar{t} \approx 0.993$, range 0.991–0.995 across reps) absent from aarch64 and riscv64 (KW $H \in \{667\,664, 667\,777, 667\,881\}$ per replication, all $p < 2.2 \times 10^{-16}$). The residual is attractor-neutral and reproducible across seeds.
4. **Algorithmic determinism across ISAs demonstrated.** The original aarch64 and riscv64 runs produced *byte-for-byte identical* CSV output across all 261 095 ticks — two architectures with distinct register files, memory models, and interrupt controllers converging to identical observable physics. Subsequent clean re-runs confirm the attractor is reproducible even when heat trajectories diverge, establishing that the SSM converges to the same fixed point from multiple initial thermal conditions.
5. **LithosAnanke boots successfully on all three ISAs.** All three QEMU acceptance runs completed POST (936+ tests), reached the ok> REPL prompt, and emitted valid DoE CSV data.
6. **Compudynamic feedback mitigates ACL enforcement overhead.** The word-level ACL capsule (ACL.4th) introduced a static-floor TTL that caused a 800× ISA overhead gap (+0.08 % on AMD64 vs. +62–68 % on AARCH64/RISCV64) when ACL-BASE-TTL = 16. Applying the same ring-buffer and slope-inference topology used by Loops #2+#3+#6 to the per-word ACL TTL produces the ACL Rolling Window of Truth (ACL-RWT), which closes the ISA gap to <2× at sub-0.01 % absolute overhead across all three architectures.
7. **ACL-RWT preserves compudynamic invariance.** All nine cells of the ACL-RWT 3 × 3 Latin-square campaign record +0.0054–+0.0088 % overhead with 0.000 % coefficient of variation in execution rate. The SSM absorbs the ACL cost into its normal feedback without perturbing the breathing cycle or the execution-rate null result. Security enforcement and physics-driven scheduling are provably non-interfering.

These results establish the SSM as a genuinely portable computational-physics engine. The compudynamic attractor is a feature of the algorithm, not of any particular silicon — and the same compudynamic principle that drives execution scheduling also governs security enforcement at negligible cost.

A Reproducibility

A.1 Build

```
git clone https://github.com/rajames440/starforth
git checkout lithosananke
make -f Makefile.starkernel ARCH=amd64 clean qemu DOE_REPS=30
make -f Makefile.starkernel ARCH=aarch64 clean qemu DOE_REPS=30
make -f Makefile.starkernel ARCH=riscv64 clean qemu DOE_REPS=30
```

CSV output appears on the QEMU serial stream and is captured automatically to `experiments/bare_metal/runs/`.

A.2 Analysis

```
cd experiments/bare_metal/analysis
Rscript analyse_bare_metal.R      # single-rep charts (latest/ data)
Rscript analyse_multirep.R        # multi-rep analysis (all 9 cells)
python3 svg_to_pdf.py
cd report && pdflatex bare_metal_doe_report.tex
pdflatex bare_metal_doe_report.tex
```

R packages required: `dplyr`, `tidyr`, `ggplot2`, `viridis`, `svglite`, `patchwork`, `scales`. System dependency: `librsvg2-bin` (provides `rsvg-convert`).

A.3 Data Files (all 9 canonical cells)

Replication 1 — seed 12345:

- `runs/doe-amd64-20260612-110212.csv` (261 098 rows, 51m 39s)

- runs/doe-aarch64-20260612-132412.csv (261 095 rows, 5m 05s)
- runs/doe-riscv64-20260612-135612.csv (261 095 rows, 5m 16s)

Replication 2 — seed 67890:

- runs/doe-riscv64-20260612-141129.csv (261 095 rows, 5m 21s)
- runs/doe-aarch64-20260612-141722.csv (261 095 rows, 5m 10s)
- runs/doe-amd64-20260612-142323.csv (261 098 rows, ~51m)

Replication 3 — seed 13579:

- runs/doe-aarch64-20260612-160132.csv (261 095 rows, ~5m)
- runs/doe-amd64-20260612-160656.csv (261 098 rows, ~51m)
- runs/doe-riscv64-20260612-174414.csv (261 095 rows, ~5m)

All files are tracked in Git LFS under `experiments/bare_metal/runs/`. The `latest/` symlink directory holds the Rep 3 files as the current canonical single-rep dataset.

A.4 Acceptance Logs

- logs2/qemu-amd64-20260612-110212.log (Rep 1)
- logs2/qemu-aarch64-20260612-132412.log (Rep 1)
- logs2/qemu-riscv64-20260612-135612.log (Rep 1)
- logs2/qemu-riscv64-20260612-141129.log (Rep 2)
- logs2/qemu-aarch64-20260612-141722.log (Rep 2)
- logs2/qemu-amd64-20260612-142323.log (Rep 2)
- logs2/qemu-aarch64-20260612-160132.log (Rep 3)
- logs2/qemu-amd64-20260612-160656.log (Rep 3)
- logs2/qemu-riscv64-20260612-174414.log (Rep 3)

A.5 ACL-RWT Campaign Data (June 2026)

All ACL-RWT runs are on branch `feature/acl-rwt` with `ACL-BOOT-RW` as the default boot policy.

ACL Floor = 16 (post Bug-1 fix, 15 June 2026):

- runs/doe-amd64-20260615-103843.csv (261 307 ticks)
- runs/doe-amd64-20260615-142839.csv (261 307 ticks)
- runs/doe-amd64-20260615-161023.csv (261 307 ticks)
- runs/doe-aarch64-20260615-121103.csv (422 431 ticks)
- runs/doe-aarch64-20260615-123814.csv (432 403 ticks)
- runs/doe-aarch64-20260615-160204.csv (432 111 ticks)
- runs/doe-riscv64-20260615-122026.csv (433 210 ticks)
- runs/doe-riscv64-20260615-122930.csv (436 040 ticks)
- runs/doe-riscv64-20260615-174420.csv (438 938 ticks)

ACL Floor = 256 (partial, 15 June 2026):

- runs/doe-amd64-20260615-204632.csv (261 248 ticks)
- runs/doe-amd64-20260615-230419.csv (261 112 ticks)
- runs/doe-aarch64-20260615-204232.csv (271 443 ticks)
- runs/doe-riscv64-20260615-214335.csv (272 725 ticks)

ACL-RWT: Replication 1 — seed 12345:

- runs/doe-amd64-20260615-103843.csv (261 113 ticks, +0.005 %)
- runs/doe-aarch64-20260615-121103.csv (261 118 ticks, +0.008 %)
- runs/doe-riscv64-20260615-122026.csv (261 118 ticks, +0.008 %)

ACL-RWT: Replication 2 — seed 67890:

- runs/doe-riscv64-20260616-000114.csv (261 117 ticks, +0.008 %)
- runs/doe-aarch64-20260616-000920.csv (261 118 ticks, +0.008 %)
- runs/doe-amd64-20260616-001614.csv (261 112 ticks, +0.005 %)

ACL-RWT: Replication 3 — seed 13579:

- runs/doe-aarch64-20260616-011049.csv (261 117 ticks, +0.008 %)
- runs/doe-amd64-20260616-011447.csv (261 113 ticks, +0.005 %)
- runs/doe-riscv64-20260616-021123.csv (261 117 ticks, +0.008 %)

A.6 ACL Build Instructions

```
git checkout feature/acl-rwt
# Baseline (no ACL):
make -f Makefile.starkernel ARCH=amd64 clean qemu DOE_REPS=30
# ACL-RWT active (default on this branch):
make -f Makefile.starkernel ARCH=amd64 clean qemu DOE_SEED=12345 DOE_REPS=30
make -f Makefile.starkernel ARCH=aarch64 clean qemu DOE_SEED=12345 DOE_REPS=30
make -f Makefile.starkernel ARCH=riscv64 clean qemu DOE_SEED=12345 DOE_REPS=30
```

To compile this report (requires texlive with pgfplots):

```
cd experiments/bare_metal/analysis/report
pdflatex bare_metal_doe_report.tex
pdflatex bare_metal_doe_report.tex # second pass for TOC
```