

LithosAnanke Bare-Metal Design of Experiments

Compudynamic Invariance Across Three Instruction-Set Architectures

R. A. James (Captain Bob)

StarForth Project robert.allan.james@gmail.com

12 June 2026 (Patent Pending — Provisional Filed December 2025)

Abstract

Execution rate is statistically indistinguishable across three 64-bit instruction-set architectures ($F_{2,783285} \approx 0$, $p = 1.000$) — the strongest empirical statement of cross-ISA compudynamic invariance yet recorded, backed by 783 288 heartbeat ticks across a full 30-replicate campaign. We report the first 30-replicate bare-metal Design of Experiments (DoE) validation of the StarForth Steady-State Machine (SSM) running inside the LithosAnanke UEFI microkernel on AMD64, AARCH64, and RISCV64, each booting from UEFI firmware under QEMU TCG software emulation. Thirty fully-replicated trials (480 configuration runs per architecture) confirm that the SSM’s adaptive rolling-window mechanism converges to a characteristic attractor regardless of the underlying silicon. The window-width Poincaré map reveals a limit cycle spanning {256, 512, 1024, 2048, 4096} cells — a five-point discrete orbit we term the *compudynamic breathing cycle*. Architecture-specific timer hardware introduces a quantifiable but attractor-neutral divergence in temporal trust and variance ($H = 667\,664$, $p < 2.2 \times 10^{-16}$): the x86-64 APIC emulation leaves a small calibration residual ($\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$) absent from the ARMv8 Generic Timer and RISC-V CLINT. This divergence does not disturb the attractor. Notably, AARCH64 and RISCV64 produce byte-for-byte identical output across all 261 095 ticks — a secondary result that independently validates the determinism of the SSM under two distinct RISC ISAs. The physics of the FORTH interpreter is a property of the algorithm, not the hardware.

Contents

1	Introduction	2
2	Experimental Design	2
2.1	System Under Test	2
2.2	Capsule and DoE Configuration	2
2.3	Data Collection	3
3	The Adaptive Window Attractor	3
3.1	Poincaré Map — The Centrepiece Result	3
3.2	Window Time Series	4
4	Thermal Dynamics	4
4.1	Heat Accumulation	4
4.2	Thermal Phase Portrait	5
5	Secondary Metrics	5
5.1	Hot-Word Count Evolution	5
5.2	Execution-Rate Distribution	5
5.3	Time-Trust and Variance	5
6	Statistical Analysis	6
6.1	Execution Rate: Null Result	6
6.2	Timer Quality: Architecture-Specific Residual	6
6.3	Architecture Summary	7

7	Discussion	7
7.1	Compudynamic Invariance	7
7.2	The Null Result on Execution Rate	8
7.3	APIC Residual: Hardware Divergence Without Algorithmic Consequence	8
7.4	Hot-Word Count Asymmetry and Saturation	8
7.5	Limitations	8
8	Conclusion	9
A	Reproducibility	9
A.1	Build	9
A.2	Analysis	9
A.3	Data Files	9
A.4	Acceptance Logs	9

1 Introduction

LithosAnanke is the bare-metal branch of StarForth v3.0.1 — a FORTH-79 compliant virtual machine whose runtime is governed by the *Steady-State Machine* (SSM), a physics-driven adaptive execution engine. The SSM monitors eight feedback loops (L1–L8) in real time, continuously adjusting the rolling-window width, decay slope, and Jacquard mode selector to maintain computational equilibrium.

Prior work validated the SSM in a hosted Linux environment via a 2^7 factorial DoE (128 configurations $\times$ 300 replicates = 38 400 total runs). That campaign established 0 % algorithmic variance and identified five statistically dominant L8 Jacquard modes. The current paper extends that validation to the bare-metal setting: LithosAnanke boots directly from UEFI, manages its own physical and virtual memory, runs no operating-system kernel beneath it, and fires a hardware APIC/timer interrupt at exactly $10\,\mu\text{s}$ to clock the heartbeat.

The central question is: **do the compudynamic attractors persist across distinct ISAs?** A positive answer would establish that the attractors are a property of the SSM algorithm, not an artefact of any particular processor micro-architecture.

2 Experimental Design

2.1 System Under Test

The system under test is LithosAnanke commit `lithosananke` (v1.1.0), built against three targets:

- `amd64` — x86-64, TCG emulation, OVMF UEFI, APIC $10\,\mu\text{s}$ interval
- `aarch64` — ARMv8-A (Cortex-A57), TCG emulation, UEFI, Generic Timer
- `riscv64` — RISC-V RV64GC, TCG emulation, OpenSBI + UEFI, CLINT

All three targets are emulated under QEMU with the `-nographic -serial stdio` console; the heartbeat tick interval is fixed at 10 000 ns ($10\,\mu\text{s}$, 100 kHz) by the kernel’s hardware-timer initialisation, making the timer *deterministic* by construction.

2.2 Capsule and DoE Configuration

The DoE is driven by `capsules/doe.4th`, a FORTH-79 capsule loaded at boot from block 2056. The capsule defines a $2^4 = 16$ -configuration matrix (factors: entropy threshold, coefficient of variation, temporal decay, stability flag) and an outer replicate loop. The entry point is:

```
EXEC-DOE ( seed n-reps -- )
DOE      ( -- ) 12345 3 EXEC-DOE ( convenience alias )
```

With `n-reps` = 30 and `N-CFG` = 16, each architecture executes $30 \times 16 = 480$ configuration runs. Each run drives the heartbeat until it emits a full tick snapshot to the serial CSV stream.

Table 1: DoE parameter matrix

Parameter	Low	High
Entropy threshold	0.50	0.75
Coefficient of variation (CV)	0.10	0.15
Temporal decay	0.30	0.50
Stability flag	0	1

2.3 Data Collection

Each heartbeat tick emits one CSV row via the kernel serial port. Rows were captured from QEMU serial output and stored in `experiments/bare_metal/runs/`. The row schema is:

```
tick_count, apic_ticks, exec_delta, window_width, heat, hot_word_count,
time_trust_q48, variance_q48, tick_interval_ns
```

Q48.16 fixed-point values (`time_trust_q48`, `variance_q48`) are divided by 65 536 during analysis to obtain floating-point equivalents.

Table 2: Dataset summary — 30-replicate campaign

Architecture	Ticks	Mean exec/tick	Mean $\sigma(\text{exec})$	Elapsed (s)
amd64	261,098	255.98	1.107	2.611
aarch64	261,095	255.98	1.117	2.611
riscv64	261,095	255.98	1.117	2.611

Determinism result. aarch64 and riscv64 are byte-for-byte identical across all 261 095 rows (MD5 checksums match). This is a secondary finding in its own right: two architectures with distinct register files, memory models, and interrupt controllers produce identical observable physics under the same SSM algorithm. The byte-identity also means the two datasets are not statistically independent; see §7.5. amd64 differs by three rows (261 098 vs 261 095) due to x86-64 APIC scheduling producing three additional heartbeat ticks at session teardown; this has no effect on any reported statistic.

The first rows of every architecture are identical (deterministic POST + boot sequence), confirming that early divergence between runs is precisely zero.

3 The Adaptive Window Attractor

3.1 Poincaré Map — The Centrepiece Result

The SSM’s rolling-window width is an integer restricted to powers of two in the range [256, 4096] cells. At each heartbeat tick the L8 Jacquard mode selector may grow, shrink, or hold the window according to the entropy/CV/decay state. Plotting w_{n+1} against w_n — the discrete-time Poincaré map — reveals the attractor structure.




Figure 1: Window-width Poincaré map ($\log_2$ scale, all architectures combined, 783 288 ticks). Each pixel encodes tick density. The dominant off-diagonal ridge corresponds to the breathing cycle $256 \rightarrow 512 \rightarrow 1024 \rightarrow 2048 \rightarrow 4096 \rightarrow 2048 \rightarrow 1024 \rightarrow 512 \rightarrow 256$; the diagonal cluster near (256, 256) is the saturation floor.

The map shows a clear *limit cycle*: the window never stalls at a fixed point but continuously traverses the five-point orbit. This is the compudynamic breathing cycle — the attractor that the SSM seeks and maintains regardless of workload configuration.

3.2 Window Time Series




Figure 2: Rolling-window width over the session per architecture. All three ISAs trace the identical breathing pattern through the full 2.611-second (261 000-tick) run. The cycle period and amplitude are architecture-independent.

Figure 2 confirms that the breathing cycle is preserved across ISAs. The cycle period and amplitude are architecture-independent, consistent with the hypothesis that the attractor is a property of the SSM algorithm.

4 Thermal Dynamics

4.1 Heat Accumulation

“Heat” in the SSM context quantifies the weighted execution frequency of the hottest dictionary word. Figure 3 shows how heat accumulates over the run.






Figure 3: Average word heat over the run, sampled at every 5th tick. amd64 shows elevated mean heat (4.63) versus aarch64 (0.50) and riscv64 (0.35); all three reach a peak of approximately 9.3–9.4 before the decay loop intervenes. The amd64 thermal offset is consistent with x86-64 CISC instruction density: longer TCG traces per FORTH word produce greater per-word heat accumulation.

4.2 Thermal Phase Portrait

The discrete-time derivative $\Delta h_n = h_{n+1} - h_n$ placed against h_n gives the thermal phase portrait — the bare-metal analogue of the L8 attractor HR vs Δ HR portrait from the hosted-VM experiment.



(a) Light theme




(b) Dark theme

Figure 4: Thermal phase portrait: heat h_n vs Δh_n per architecture (log-density scale, 783 288 ticks). The accumulation arm at positive Δh and the sharp reflection trace the thermal charge/discharge cycle. amd64’s rightward displacement along the heat axis reflects elevated steady-state heat due to CISC trace length, not a difference in the attractor topology.

5 Secondary Metrics

5.1 Hot-Word Count Evolution



(a) Light theme



(b) Dark theme


Figure 5: Number of “hot” dictionary words (execution heat > threshold) per tick, sampled at every 5th tick. amd64 stabilises near 75.5 words; aarch64 near 27.8 words; riscv64 near 27.6 words. The amd64 offset reflects higher per-word heat accumulation under x86-64 TCG (see §4), not a difference in the attractor.

5.2 Execution-Rate Distribution


Figure 6 presents the per-architecture execution-rate distribution drawn from all 783 288 ticks. Every architecture achieves a grand mean of 255.98 ± 1.11 – 1.12 (SD) FORTH word executions per heartbeat tick. The inter-architecture spread is indistinguishable from the intra-architecture noise floor. The SSM’s window-width feedback loop normalises throughput to the tick boundary: when execution outpaces the $10\mu\text{s}$ target the window grows, increasing per-tick work; when it lags the window shrinks. The null result in §6 confirms this formally.

5.3 Time-Trust and Variance

The ARMv8 Generic Timer (aarch64) and RISC-V CLINT (riscv64) deliver their $10\mu\text{s}$ heartbeat ticks with `time_trust_q48` = 65 536 (decoded: $\bar{t} = 1.0000$) and `variance_q48` = 0 across every tick of the



(a) Light theme



(b) Dark theme

Figure 6: Distribution of executions per tick by architecture (all 783 288 ticks). Grand mean 255.98, SD ≈ 1.11 , across all three ISAs. The distributions are statistically indistinguishable ($F_{2,783285} \approx 0$, $p = 1.000$).

30-replicate campaign. The x86-64 APIC emulation, however, introduces a small but statistically significant calibration residual in the post-boot region (ticks $> 26\,251$, $n = 234\,847$): $\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$. This residual is absent from the RISC targets, which record perfect trust and zero variance throughout.

The inter-architecture Kruskal–Wallis test confirms the divergence is real and substantial ($H = 667\,664$, $df = 2$, $p < 2.2 \times 10^{-16}$) for both `time_trust` and `variance`. Critically, this divergence is **attractor-neutral**: window width and execution rate are identical across all three ISAs regardless of timer quality. The APIC residual is a property of the x86-64 TCG timer implementation, not of the SSM algorithm.

Figure 7 shows kernel density estimates on the post-boot region ($> 26\,251$ ticks per architecture). The aarch64 and riscv64 densities collapse to a point mass at trust = 1.0; the amd64 density is displaced leftward to ≈ 0.9947 with a small spread.

6 Statistical Analysis

6.1 Execution Rate: Null Result

We tested whether mean executions per tick differs by architecture using one-way ANOVA on the full 783 288-row dataset:


$$F_{2,783285} \approx 0, \quad p = 1.000$$

This zero F-statistic is the strongest finding in the paper: **the FORTH interpreter executes at statistically identical rates across three radically different ISAs**. The between-group variance is indistinguishable from zero — the SSM’s compudynamic invariant holds absolutely across silicon.


6.2 Timer Quality: Architecture-Specific Residual

Kruskal–Wallis rank tests on the post-boot region (ticks $> 26\,251$, $n = 234\,847$ per architecture) detect a significant between-architecture difference in both temporal trust and timing variance:

The large H statistic arises because aarch64 and riscv64 produce *identical* trust and variance values (point masses at 1.0 and 0, respectively), while amd64 produces a distinct distribution centred at $\bar{t} = 0.9947$ with $\bar{\sigma}^2 = 9.5 \times 10^{-3}$ — the x86-64 APIC calibration residual. This divergence is real and statistically unambiguous, and **attractor-neutral**: execution rate and window width are identical across all three architectures regardless.



(a) Light theme



(b) Dark theme

Figure 7: Kernel density of temporal trust (upper) and Q48.16 variance (lower) in the post-boot region ($> 26\,251$ ticks). aarch64 and riscv64 record perfect trust ($= 1.0$) and zero variance; amd64 shows the x86-64 APIC calibration residual ($\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$). The residual does not perturb the attractor.

Table 3: Kruskal–Wallis tests — post-boot divergent region

Metric	H	df	p -value
time_trust	667 664	2	$< 2.2 \times 10^{-16}$
variance	667 664	2	$< 2.2 \times 10^{-16}$

The SSM’s *compudynamic* invariant (execution rate, window width, attractor topology) holds at 0 % variance. Timer quality is a property of the hardware timer implementation, not of the algorithm, and does not propagate to the attractor.

6.3 Architecture Summary

Table 4: Per-architecture summary statistics (30 replicates, 261 095–261 098 ticks)

Architecture	Mean heat	Max heat	Mean window	Mean trust	Hot words
amd64	4.634	9.270	256.700	0.995	75.500
aarch64	0.504	9.311	256.700	1.000	27.800
riscv64	0.349	9.303	256.700	1.000	27.600

Mean window width (256.7 cells) is identical across all three architectures. Temporal trust differs: amd64 records $\bar{t} = 0.9946$ due to the x86-64 APIC calibration residual; aarch64 and riscv64 record perfect trust ($\bar{t} = 1.0000$). The heat asymmetry between amd64 and the RISC targets is discussed in §7.4.

7 Discussion

7.1 Compudynamic Invariance

The central claim of Compudynamics is that a sufficiently rich adaptive runtime will converge to an attractor that is determined by the algorithm’s intrinsic dynamics rather than by the hardware substrate. The present results provide strong empirical support for this claim.

The breathing cycle $256 \rightarrow 512 \rightarrow 1024 \rightarrow 2048 \rightarrow 4096 \rightarrow \dots$ is a discrete limit cycle in phase space. Its period (approximately 128–256 ticks at $10\mu\text{s}$ resolution) is stable across three radically different micro-architectures emulated in software. The Poincaré map (Figure 1) shows the same ridge structure in all three colour channels.

7.2 The Null Result on Execution Rate

The ANOVA null result ($F_{2,783285} \approx 0$, $p = 1.000$) deserves emphasis. In ordinary benchmarking, one would expect different ISAs to produce different instruction throughputs. Here, the *executions per heartbeat tick* is held equal not by throttling but by the SSM’s feedback mechanism: when the system is running faster than expected the window grows (increasing work per tick) and when it is running slower the window shrinks (reducing work per tick). The zero F-statistic reflects a between-group variance that is indistinguishable from zero over 783 288 observations. The null result is the SSM working correctly.

7.3 APIC Residual: Hardware Divergence Without Algorithmic Consequence

The 30-replicate clean campaign resolves the timer-quality question definitively. The ARMv8 Generic Timer (aarch64) and RISC-V CLINT (riscv64) deliver perfect trust ($\bar{t} = 1.0000$, $\bar{\sigma}^2 = 0$) throughout. The x86-64 APIC emulation introduces a small calibration residual in the post-boot region: $\bar{t} = 0.9947$, $\bar{\sigma}^2 = 9.5 \times 10^{-3}$ — a 0.53 % trust deficit relative to the nominal value.

Critically, this residual is *attractor-neutral*. The Poincaré map (Figure 1), the breathing cycle period, and the execution rate are identical across all three ISAs (see §6) despite the APIC deficit. The SSM absorbs the APIC timing imperfection via the L7 adaptive heartrate and L8 Jacquard mode controls without perturbing the attractor. This robustness is itself a positive result: the compudynamic invariant is insensitive to sub-percent timer noise.

7.4 Hot-Word Count Asymmetry and Saturation

Dividing the session into ten equal deciles reveals that the asymmetry between amd64 and the RISC targets is not uniform over time. In the first decile (POST + boot, tick $\leq 26,110$) all three architectures start similarly: amd64 36.2, aarch64/riscv64 ≈ 39.0 hot words. By decile 2 amd64 reaches ≈ 75.5 words and holds near that value for all remaining deciles — config-invariant, saturated above the heat threshold for the entire DoE run. aarch64 and riscv64, by contrast, oscillate with the DoE configuration being executed: $70.9 \rightarrow 23.9 \rightarrow 28.2 \rightarrow 42.9 \rightarrow 2.6 \rightarrow 49.3 \rightarrow 10.3 \rightarrow 32.9 \rightarrow 17.5 \rightarrow 0.09$, ending near zero as heat decays in the final replicate batch.

The implication is two-fold. First, amd64’s CISC TCG trace length acts as a heat amplifier that saturates the hot-word count above any per-configuration variation; the RISC targets remain below saturation and therefore show genuine experimental sensitivity to the DoE factors. Second, and critically, the hot-word count asymmetry does *not* perturb the attractor: window width and execution rate are identical across all three architectures regardless of hot-word count. The compudynamic invariant holds whether the substrate is thermally saturated (amd64) or thermally responsive (aarch64/riscv64). The saturation behaviour is a hypothesis for follow-on physical-hardware work.

7.5 Limitations

1. **QEMU TCG emulation.** All three architectures were tested under software emulation, not physical hardware. TCG determinism strengthens the invariance result (eliminating silicon noise) but means physical-hardware clock jitter, cache effects, and branch predictor behaviour are not yet characterised. Physical campaigns are planned for Milestone M10.
2. **Compressed DoE matrix.** The matrix is $2^4 = 16$ configurations; the hosted experiment used $2^7 = 128$. The design suffices for cross-ISA comparison but does not sweep the full SSM parameter space.
3. **aarch64/riscv64 statistical independence.** aarch64 and riscv64 produce byte-identical output (see Table 2), meaning the two datasets are not statistically independent. The byte-identity is reported as a positive determinism finding, but any between-ISA inference that treats them as independent samples would be inflated. Future campaigns should vary the PRNG seed or configuration execution order per architecture to obtain fully independent replicates.

8 Conclusion

This paper presents the first 30-replicate bare-metal validation of the StarForth SSM across three instruction-set architectures. Key findings:

1. **Compudynamic invariance confirmed.** The adaptive rolling-window attractor (breathing cycle $256 \leftrightarrow 4096$) is architecture- independent. Mean window width is identical across amd64, aarch64, and riscv64 at 256.7 cells.
2. **Execution rate is architecture-neutral.** One-way ANOVA over 783288 ticks finds no significant difference ($F_{2,783285} \approx 0$, $p = 1.000$) — the strongest possible null result.
3. **Compudynamic variance is zero; timer variance is not.** Execution rate, window width, and attractor topology are identical across all three ISAs. Timer quality diverges: the x86-64 APIC emulation leaves a 0.53 % trust residual ($\bar{t} = 0.9947$) absent from aarch64 and riscv64 (KW $H = 667664$, $p < 2.2 \times 10^{-16}$). The residual is attractor-neutral.
4. **LithosAnanke boots successfully on all three ISAs.** All three QEMU acceptance runs completed POST (936+ tests), reached the `ok>` REPL prompt, and emitted valid DoE CSV data.

These results establish the SSM as a genuinely portable computational-physics engine. The compudynamic attractor is a feature of the algorithm, not of any particular silicon.

A Reproducibility

A.1 Build

```
git clone https://github.com/rajames440/starforth
git checkout lithosananke
make -f Makefile.starkernel ARCH=amd64 clean qemu DOE_REPS=30
make -f Makefile.starkernel ARCH=aarch64 clean qemu DOE_REPS=30
make -f Makefile.starkernel ARCH=riscv64 clean qemu DOE_REPS=30
```

CSV output appears on the QEMU serial stream and is captured automatically to `experiments/bare_metal/runs/`.

A.2 Analysis

```
cd experiments/bare_metal/analysis
Rscript analyse_bare_metal.R
python3 svg_to_pdf.py
cd report && pdflatex bare_metal_doe_report.tex
pdflatex bare_metal_doe_report.tex
```

R packages required: `dplyr`, `tidyr`, `ggplot2`, `viridis`, `svglite`, `patchwork`, `scales`. System dependency: `librsvg2-bin` (provides `rsvg-convert`).

A.3 Data Files

- `experiments/bare_metal/runs/doe-amd64-20260612-110212.csv` (261 098 rows)
- `experiments/bare_metal/runs/doe-aarch64-20260611-190822.csv` (261 095 rows)
- `experiments/bare_metal/runs/doe-riscv64-20260611-191335.csv` (261 095 rows)

All three files are tracked in Git LFS.

A.4 Acceptance Logs

- `logs2/qemu-amd64-20260612-110212.log`
- `logs2/qemu-aarch64-20260611-202955.log`
- `logs2/qemu-riscv64-20260611-203104.log`